

Access Code: **318159**

CMP9794M

Advanced Artificial Intelligence

[Heriberto Cuayahuitl](#)



UNIVERSITY OF
LINCOLN

School of Engineering and Physical Sciences

Last Week

- Introduction to Bayesian Networks
 - Networks encode conditional independence
 - Structure via a Directed Acyclic Graph (DAG)
 - Each node has a conditional prob. table (CPT)
 - CPTs can be learnt via Maximum Likelihood Estimation (MLE) with Laplace / Aditive / Dirichlet smooting to avoid zero probabilities
- Introduction to Probabilistic Reasoning
 - Inference by enumeration
 - Inference by variable elimination

Today: Structure Learning

- It is often difficult, if not impossible, to manually specify the structure of a Bayesian network.
 - This difficulty increases according to the number of random variables involved.
 - Due to the lack of agreement between domain experts
- Bayesian network learning is defined as

$$\underbrace{\Pr(\mathcal{B}|\mathcal{D})}_{\text{Bayesian network}} = \underbrace{\Pr(G, \Theta|\mathcal{D})}_{\text{Data}} = \underbrace{\Pr(G|\mathcal{D})}_{\text{structure learning}} \cdot \underbrace{\Pr(\Theta|G, \mathcal{D})}_{\text{parameter learning}} .$$

learning

Directed Acyclic Graph (DAG)

Parameters (probability distributions)

Complexity of Structure Learning

Table 1 Number of directed graphs, directed acyclic graphs, and the percentage of directed graphs which are acyclic for different number of variables


Number of variables, n	Number of directed graphs $(3^{n(n-1)/2})$	Number of DAGs, $ G_n $	Percentage of directed graphs which are acyclic (%)
2	3	3	100.0
3	27	25	92.59
4	729	543	74.49
5	59,049	29,281	49.59
6	1.4349×10^7	3.7815×10^6	26.35
7	1.0460×10^{10}	1.1388×10^9	10.89
8	2.2877×10^{13}	7.8730×10^{11}	3.42

This calculation assumes that the directed graph has at most one arc between each pair of nodes

Today

- **Constraint-based structure learning**
 - PC-stable algorithm
- Score-based structure learning
 - Greedy search (Hill Climbing)
- Hybrid algorithms
 - Min-Max Hill Climbing
- Evaluation metrics for probabilistic models

PC-Stable Algorithm: Key Steps



PC-Stable: Find Skeleton

Algorithm 1. The First Step in PC-Stable Algorithm

Input: \mathcal{V}

Output: $G, SepSet$

```
1:  $G =$  fully connected graph
2:  $SepSet = \emptyset$ 
3:  $\ell = 0$ 
4: repeat
5:   Copy  $G$  into  $G'$ 
6:   for any edge  $(V_i, V_j)$  in  $G$  do
7:     repeat
8:       Choose a new  $S \subseteq adj(V_i, G') \setminus \{V_j\}$  with  $|S| = \ell$ 
9:       Perform  $I(V_i, V_j | S)$ 
10:      if  $V_i \perp\!\!\!\perp V_j | S$  then
11:        Remove  $(V_i, V_j)$  from  $G$ 
12:        Store  $S$  in  $SepSet$ 
13:      end if
14:    until  $(V_i, V_j)$  is removed or all sets  $S$  are considered
15:  end for
16:   $\ell = \ell + 1$ 
17: until ( max degree  $- 1 \geq \ell$  )
```

Knowing variables S, V_i
does not contribute to
knowing more about V_j .

PC-Stable: Example Skeleton

			time ↑
		I (3,0 {2}) = × I (3,0 {1}) = ×	
	G		I (2,3 {0}) = ✓ I (2,0 {3}) = ×
	I (2,3) = × I (1,3) = ×		I (1,3 {0}) = ✓ I (1,0 {3}) = ×
	G		I (0,3 {2}) = × I (0,3 {1}) = ×
	I (1,2) = ✓ I (0,3) = × I (0,2) = × I (0,1) = ×		I (0,2 {3}) = × I (0,2 {1}) = × I (0,1 {3}) = × I (0,1 {2}) = ×
	G		I (0,1 {2,3}) = × I (0,2 {1,3}) = × I (0,3 {1,2}) = ×
Level 0		 Create complete graph G	 Copy G into G'
			 Copy G into G'
			Level 2

Conditional Independence Tests

Assuming that each arc (in a Bayesian network) encodes a probabilistic dependence, conditional independence tests tell us whether such probabilistic dependence is supported by the data.

Any two variables are associated in a Bayes net if the test's p-value < significance_level.

→ typically **0.05** or **0.01**
(default=0.05)

If $p\text{-value} < \text{significance level}$: **keep edges**
Else: **remove edges**

API for Conditional Independence Tests

```
bnlearn.bnlearn.independence_test(model, data  
, test='chi_square', alpha=0.05, prune=False,  
verbose=3)
```

The statistical tests:

- chi_square
- g_sq
- log_likelihood
- freeman_tuckey
- modified_log_likelihood
- neyman
- cressie_read

Look for the following in
this week's workshop:
`bnlearn_ConditionalIndependenceTests.py`

API for PC-Stable Algorithm


```
bnlearn.structure_learning.fit(data, methodt  
ype='pc', params_pc={'alpha': 0.05,  
'ci_test': 'chi_square'}, verbose=3)
```

The statistical tests:

- chi_square
- g_sq
- log_likelihood
- freeman_tuckey
- neyman
- cressie_read
- power_divergence
- pearsonr

Look for the following in
this week's workshop:
`bnlearn_ConditionalIn
dependenceTests.py`


Conditional Independence Tests using the Language Detection data



Knowing Y: Is the link between X_1 and X_2 with a p-value of 0.0 needed?

Is the link between X_1 and X_{15} with a p-value of 0.4692 needed?

Conditional Independence Tests using the Language Detection data



Knowing Y: Is the link between X_1 and X_2 with a p-value of 0.0 needed?
Yes, keep the link.

Is the link between X_1 and X_{15} with a p-value of 0.4692 needed?
No, remove the link.

Today

- Constraint-based structure learning
 - PC-stable algorithm
- **Score-based structure learning**
 - Greedy search (Hill Climbing)
- Hybrid algorithms
 - Min-Max Hill Climbing
- Evaluation metrics for probabilistic models

Greedy Search: Main Ideas

Explore the search space starting from an empty network and **adding**/**deleting**/**reversing** one arc at a time until the score can't be improved anymore.

Doing the above requires knowledge regarding how good a network is against alternative graphs.

The goal is to search for the network with the *highest score*: $B^* = \arg \max_B Score(B|D)$

Bayesian network Data

Scoring Function

The score represents a trade-off:

- How well the network fits the data &
 - How complex the network is

Assumption: the scoring function is decomposable

$$Score(B|D) = \sum_{i=1}^n Score(X_i | parents(X_i), D)$$

random variable i

data

Several scoring functions can be applied to structure learning such as LL, BIC, AIC, BDeu, etc.

Scoring Functions: Log-Likelihood (LL)

LL represents the probability of the data (D) given a Bayesian network (B):

$$\begin{aligned} LL(D|B) &= \sum_j^N \log P(D_j|B) \\ &= \sum_i^n \sum_j^N \log P(D_{ij}|PA_{ij}) \end{aligned}$$

where D_{ij} is the value of random variable X_i in training example D_j , and PA_{ij} are the values of the parents of X_i in D_j , n is the number of random variables, and N is the num. of training examples.

LL Example for PlayTennis Data using a Naïve Bayes Structure

$$LL = \sum_{i=1}^N \log P(PlayTennis_i) + \sum_i^n \sum_j^N \log P(X_{ij}|PA_i)$$

$$\begin{aligned} LL_1 &= \log P(PlayTennis = no) \\ &+ \log P(Outlook = sunny|PlayTennis = no) + \\ &+ \log P(Temperature = hot|PlayTennis = no) + \\ &+ \log P(Humidity = high|PlayTennis = no) + \\ &+ \log P(Wind = weak|PlayTennis = no) \\ &= \log\left(\frac{5}{14}\right) + \log\left(\frac{3}{5}\right) + \log\left(\frac{2}{5}\right) + \log\left(\frac{4}{5}\right) + \log\left(\frac{2}{5}\right) = -3.5961 \end{aligned}$$

and so on for $LL_1 + \dots + LL_{14} = -54.217$

Scoring Functions: Bayesian Information Criterion (BIC)

BIC is a penalised LL function:

$$BIC(D|B) = LL(D|B) - \sum_{i=1}^n Penalty(X_i, B, D)$$

where $LL(\cdot)$ is Log-Likelihood function and the penalty term penalises complex networks as

$Penalty(X_i, B, D) = \frac{\log N * p_i}{2}$, p_i is the number of parameters/probabilities of random variable X_i , and N is the number of training examples.

BIC Example for PlayTennis Data using a Naïve Bayes Structure

$$\begin{aligned} BIC(D|B) &= LL(D|B) - \sum_i \frac{\log N * p_i}{2} \\ &= LL(D|B) - \frac{\log N * \sum_i p_i}{2} \end{aligned}$$

$$\begin{aligned} BIC(D|B) &= -54.217 - \frac{\log(14) * 22}{2} \\ &= -83.246 \end{aligned}$$

N.B. The above assume CPTs with full enumeration (not concise ones)

Scoring Functions: Aikake Information Criterion (AIC)

AIC is a penalised LL function:

$$AIC(D|B) = LL(D|B) - \sum_{i=1}^n Penalty(X_i, B, D)$$

where $LL(\cdot)$ is Log-Likelihood function and $Penalty(X_i, B, D) = p_i$ with p_i being the number of parameters/probabilities of random variable X_i , and N is the number of training examples.

AIC favours more complex networks than BIC.

AIC Example for PlayTennis Data using a Naïve Bayes Structure

$$AIC(D|B) = LL(D|B) - \sum_{i=1}^N p_i$$

$$\begin{aligned} AIC(D|B) &= -54.217 - (2 + 6 + 6 + 4 + 4) \\ &= -76.217 \end{aligned}$$

Using full enumeration ->

$$\begin{aligned} |CPT(H|PT)| &= 4 \\ |CPT(W|PT)| &= 4 \\ |CPT(T|PT)| &= 6 \\ |CPT(O|PT)| &= 6 \\ |CPT(PT)| &= 2 \end{aligned}$$

AIC Example for PlayTennis Data using a Naïve Bayes Structure

$$AIC(D|B) = LL(D|B) - \sum_{i=1}^N p_i$$

$$\begin{aligned} AIC(D|B) &= -54.217 - (1 + 3 + 3 + 2 + 2) \\ &= -65.217 \end{aligned}$$

Using concise enumeration ->

$$\begin{aligned} |CPT(H|PT)| &= 2 \\ |CPT(W|PT)| &= 2 \\ |CPT(T|PT)| &= 3 \\ |CPT(O|PT)| &= 3 \\ |CPT(PT)| &= 1 \end{aligned}$$


API for Structure Learning with Scoring Functions

```
bnlearn.structure_learning.fit(data,  
methodtype='hc', scoretype='bic',  
black_list=None, white_list=None, ...,  
max_iter=1000000, ..., n_jobs=-1, verbose=3)
```


The statistical tests:

- ‘bic’ : Bayesian information criterion
- ‘k2’ : K2 score
- ‘bdeu’ : Bayesian Dirichlet score
- ‘bds’ : Bayesian Dirichlet sparse score
- ‘aic’ : Akaike information criterion

LL/BIC/AIC for Language Detection



BIC=???
AIC=???
Bdeu=???



BIC=???
AIC=???
Bdeu=???


Which structure above fits better the data?

Hill Climbing Algorithm


1. Choose a network structure G over \mathbf{V} , usually (but not necessarily) empty.
2. Compute the score of G , denoted as $\text{Score}_G = \text{Score}(G)$.
3. Set $\text{maxscore} = \text{Score}_G$.
4. Repeat the following steps as long as maxscore increases:
 - (a) for every possible arc addition, deletion or reversal not resulting in a cyclic network:
 - i. compute the score of the modified network G^* , $\text{Score}_{G^*} = \text{Score}(G^*)$:
 - ii. if $\text{Score}_{G^*} > \text{Score}_G$, set $G = G^*$ and $\text{Score}_G = \text{Score}_{G^*}$.
 - (b) update maxscore with the new value of Score_G .
5. Return the DAG G .

Look for the following in this week's workshop:
bnlearn_StructureLearning.py

Depth First Search (DFS) for Detecting Cyclic Networks



DAGs do
not have
back edges



Look for function
`has_cycles()` in
BayeNetUtil.py



NOTATION:

- Tree edges
- Forward edges
- Back edges

Today


- Constraint-based structure learning
 - PC-stable algorithm
- Score-based structure learning
 - Greedy search (Hill Climbing)
- **Hybrid algorithms**
 - Min-Max Hill Climbing
- Evaluation metrics for probabilistic models

Hybrid Algorithms for Structure Learning in Bayesian Networks



- Conditional Independent tests are partially rule-based
- Score-based approaches can be search intensive
- Is there any other alternative?

Hybrid Algorithms for Structure Learning in Bayesian Networks



Since Conditional Independent tests and Score-based approaches have pros and cons, hybrid algorithms aim for **the best of both worlds**.

MMHC: Min-Max Hill Climbing

1. Choose a network structure G over V (nodes), usually empty (but not necessarily).
2. **Restrict:** select a set \mathbf{C}_i of candidate parents for each node $X_i \in V$, which must include the parents of X_i in G .
via the PC-stable algorithm
3. **Maximise:** find the network structure G^* that maximises $Score(G^*)$ among the networks in which the parents of each node X_i are included in the corresponding set \mathbf{C}_i .
via the Hill Climbing algorithm
4. Return DAG G^* .

Today

- Constraint-based structure learning
 - PC-stable algorithm
- Score-based structure learning
 - Greedy search (Hill Climbing)
- Hybrid algorithms
 - Min-Max Hill Climbing
- **Evaluation metrics for probabilistic models**
 - See appendix 1, discussion during the workshop

Today

- Discussion on structure learning
- Algorithms to induce the structure of Bayes nets
- Measuring the performance of probabilistic models/classifiers

Readings:

- [Kitson et al. A survey on Bayesian Network structure learning, arXiv, 2023](#) (main reading)
- [Liu et al. Empirical evaluation of scoring functions for Bayesian Network model selection, BMC Bioinformatics, 2012](#) (optional reading)

This and Next Week

Workshop (tomorrow):

Exercise on Conditional Independent tests

Exercise on scoring functions (BIC, AIC, BDeu)

Python code for structure learning

Lecture (next week):

Gaussian Bayesian Networks

Reading 1: Koller & Friedman 2009. [Section 7.2](#)

Reading 2: Bishop. C. 2006. [Section 8.1.4](#)

Questions?

Appendix 1

Metrics for Evaluating Probabilistic Models

Metrics for Probabilistic Models

- The higher the better:
 - Balanced accuracy, $[0, \dots, 1]$
 - F1 Score, $[0, \dots, 1]$
 - Area under curve, $[0, \dots, 1]$
- The lower the better
 - Brier score, $[0, \dots, 1]$
 - KL divergence, $[0, \dots, \infty]$
 - Expected Calibration Loss, $[0, \dots, 1]$
 - Training and inference times (in seconds)

Metrics: Balanced Accuracy

		Predicted Label/Condition	
		POSITIVE (PP)	NEGATIVE (PN)
Actual Label or Condition	POSITIVE (P)	True Positive (TP)	False Negative (FN)
	NEGATIVE (N)	False Positive (FP)	True Negative (TN)

$$BA = \frac{TPR + TNR}{2}, \text{ where}$$

$$TPR = \frac{TP}{P} = \frac{TP}{TP + FN} = 1 - FNR \quad FNR = \frac{FN}{P} = \frac{FN}{FN + TP} = 1 - TPR$$

$$TNR = \frac{TN}{N} = \frac{TN}{TN + FP} = 1 - FPR \quad FPR = \frac{FP}{N} = \frac{FP}{FP + TN} = 1 - TNR$$

If $FNR = FPR = 0$ then $TPR + TNR = 2$

Metrics: F1 Score

Actual Label or Condition	Predicted Label/Condition	
	POSITIVE (PP)	NEGATIVE (PN)
	POSITIVE (P)	True Positive (TP) False Negative (FN)
	NEGATIVE (N)	False Positive (FP) True Negative (TN)

$$F_1 = 2 * \left(\frac{Precision * Recall}{Precision + Recall} \right), \text{ where}$$

$$Precision = \frac{TP}{PP} = \frac{TP}{TP + FP}$$
 ← Of all positive predictions, what fraction is actually positive.

$$Recall = \frac{TP}{P} = \frac{TP}{TP + FN}$$
 ← Of all actual positive instances, what fraction is correctly predicted as positive.

Metrics: Area Under the ROC Curve (AUC)


Predicted Label/Condition

Actual Label or Condition	Predicted Label/Condition	
	POSITIVE (PP)	NEGATIVE (PN)
POSITIVE (P)	True Positive (TP)	False Negative (FN)
NEGATIVE (N)	False Positive (FP)	True Negative (TN)

AUC measures the entire area under the entire ROC curve – the higher the area the better the classifier.

$$TPR = \frac{TP}{P} = \frac{TP}{TP + FN} = 1 - FNR$$

$$FPR = \frac{FP}{N} = \frac{FP}{FP + TN} = 1 - TNR$$



Metrics: Area Under the ROC Curve (AUC)

Predicted Label/Condition

Actual Label or Condition		
	POSITIVE (PP)	NEGATIVE (PN)
	POSITIVE (P)	True Positive (TP) False Negative (FN)
NEGATIVE (N)	False Positive (FP)	True Negative (TN)

AUC measures the entire area under the entire ROC curve – the higher the area the better the classifier.

$$TPR = \frac{TP}{P} = \frac{TP}{TP + FN} = 1 - FNR$$

$$FPR = \frac{FP}{N} = \frac{FP}{FP + TN} = 1 - TNR$$




Figure from [this link](#)

Metrics for Probabilistic Models

- The higher the better:
 - Balanced accuracy, [0,...,1]
 - F1 Score, [0,...,1]
 - Area under curve, [0,...,1]
- The lower the better
 - Brier score, [0,...,1]
$$BS = \frac{1}{N} \sum_{t=1}^N (f_t - o_t)^2$$
 - KL divergence, [0,..., ∞]
$$D_{\text{KL}}(P \parallel Q) = \sum_{x \in \mathcal{X}} P(x) \log\left(\frac{P(x)}{Q(x)}\right)$$
 - Expected Calibration Loss, [0,...,1]
 - Training and inference times (in seconds)

KL Divergence: Example (1/2)

- True Distribution $P = <1.0, 0.0>$
- Predicted distribution $Q = <0.6, 0.4>$

$$D_{\text{KL}}(P \parallel Q) = \sum_{x \in \mathcal{X}} P(x) \log \left(\frac{P(x)}{Q(x)} \right)$$

True/Correct distribution Predicted distribution

$$\begin{aligned} D_{KL}(P||Q) &= P(x_1) \log \left(\frac{P(x_1)}{Q(x_1)} \right) + P(x_2) \log \left(\frac{P(x_2)}{Q(x_2)} \right) \\ &= 1.0 \log \left(\frac{1.0}{0.6} \right) + 0.0 \log \left(\frac{0.0}{0.4} \right) \\ &= 0.5108 + 0 \end{aligned}$$

KL Divergence: Example (2/2)

- True Distribution $P = <1.0, 0.0>$
- Predicted distribution $Q = <0.4, 0.6>$

$$D_{\text{KL}}(P \parallel Q) = \sum_{x \in \mathcal{X}} P(x) \log \left(\frac{P(x)}{Q(x)} \right)$$

$$\begin{aligned} D_{KL}(P||Q) &= P(x_1) \log \left(\frac{P(x_1)}{Q(x_1)} \right) + P(x_2) \log \left(\frac{P(x_2)}{Q(x_2)} \right) \\ &= 1.0 \log \left(\frac{1.0}{0.4} \right) + 0.0 \log \left(\frac{0.0}{0.6} \right) \\ &= 0.9162 + 0 \end{aligned}$$

N.B. This example uses the natural logarithm (base e) – default in Python. Other choices are: $\log_{10}()$ and $\log_2()$

Expected Calibration Loss (ECL)

It measures the discrepancy between predicted probabilities and expected accuracy.

$$ECL = \sum_{i=1}^N b_i \cdot |p_i - c_i|$$

High ECL means
the model is either
overconfident or
underconfident.

Low ECL means
model **well**
calibrated.

- N is the number of bins.
- b_i is the fraction of samples in the i_{th} bin.
- p_i is the avg. predicted probability in the i_{th} bin.
- c_i is the observed accuracy in the i_{th} bin.

Metrics for Probabilistic Models


- The higher the better:
 - Balanced accuracy, $[0, \dots, 1]$
 - F1 Score, $[0, \dots, 1]$
 - Area under curve, $[0, \dots, 1]$
- The lower the better:
 - Brier score, $[0, \dots, 1]$
 - KL divergence, $[0, \dots, \infty]$
 - Expected Calibration Loss, $[0, \dots, 1]$
 - Training and inference times (in seconds)

Consider using these metrics in your assignment. Look at **ModelEvaluator.py**. Fine if you wish to include any others.

Appendix 2

Evaluation of
Probabilistic Models
Using Cross Validation

K-fold Cross Validation



K-fold CV: Methodology for Pure Model Evaluation

Suitable for small datasets

1. Discretise your data, if required. ([examples](#))
2. Split the entire dataset in K folds—randomly.
3. Train on K-1 folds and test using the remaining fold.
4. Iterate over all K consistent folds.
5. Average the performance according to difference metrics, and report standard deviation if considered relevant.
6. Report results for different models in a table.